

Algorithm Design and Analysis

Day 6
Distances in Graphs

2015, AUT-CJLU

Day 6: Distances in Graphs


Part I: Breadth First Search

Traversing a Graph


Uses for Graph Traversals

- Searching (DFS)
- Reachability (DFS)
- Decomposing graphs (DFS)
- Calculating distances (DFS is not suitable)

Graph G




The DFS Tree of G



Distances Between Nodes

Recall: Let G be a graph


- The **length** of a path is the number of edges (“steps”) in it.
- The **distance** from a node u to v , $\text{dist}(u, v)$, is the length of the shortest path from u to v . If no path exist, then $\text{dist}(u, v) = \infty$.



Distances Between Nodes

Recall: Let G be a graph

- The **length** of a path is the number of edges (“steps”) in it.
- The **distance** from a node u to v , $\text{dist}(u, v)$, is the length of the shortest path from u to v . If no path exist, then $\text{dist}(u, v) = \infty$.



- Distance between D and J: 4
- Distance between A and K: ∞

Distances Between Nodes

Distance Problem

INPUT: a graph G , and two nodes u, v

OUTPUT: the distance from u to v .

Shortest Path Problem

INPUT: a graph G , and two nodes u, v

OUTPUT: the shortest path from u to v .

Case Study: Missionary and Cannibals

Missionaries and Cannibals

There are 3 missionaries and 3 cannibals coming to a river with only one boat that can hold only 2 people. At any instance the number of cannibals cannot be more than the number of missionaries. How can all the people cross the river using the least number of boat rides?



Case Study: Missionary and Cannibals

Convert the problem into a graph problem:

- We call a time stamp of the current situation a **configuration**.
- A configuration states how many missionaries and cannibals on each bank, and the location of the boat.


E.g. $(MC\bullet, MMCC)$ indicates 1 missionary + 1 cannibal on left bank, 2 missionaries + 2 cannibals on right bank, boat on left bank

Case Study: Missionary and Cannibals

Convert the problem into a graph problem:

- We call a time stamp of the current situation a **configuration**.
- A configuration states how many missionaries and cannibals on each bank, and the location of the boat.
E.g. $(MC\bullet, MMCC)$ indicates 1 missionary + 1 cannibal on left bank, 2 missionaries + 2 cannibals on right bank, boat on left bank
- Define a graph $G = (V, E)$ where
 - The nodes are all configurations
 - Two nodes are connected by an edge if from one node (configuration) the six people can move to the other node (configuration) using one boat ride.

Case Study: Missionary and Cannibals



Our goal: Find a **shortest path** from $(, \bullet MMMCCC)$ to $(MMMCCC \bullet,)$.

Breadth First Search

How to solve the distance and the shortest path problem?

Breadth First Search

How to solve the distance and the shortest path problem?

Breadth First Search Strategy

Traverse nodes by “layers”:

- ① Visit the start node s
- ② Visit all nodes that have distance 1 from s (call them V_1)
- ③ Visit all nodes that have distance 1 from V_1 (call them V_2)
- ④ Visit all nodes that have distance 1 from V_2 (call them V_3)
- ⑤



Breadth First Search

How to solve the distance and the shortest path problem?

Breadth First Search Strategy

Traverse nodes by “layers”:

- 1 Visit the start node s
- 2 Visit all nodes that have distance 1 from s (call them V_1)
- 3 Visit all nodes that have distance 1 from V_1 (call them V_2)
- 4 Visit all nodes that have distance 1 from V_2 (call them V_3)
- 5




Breadth First Search

Queue implementation of BFS

Maintain a **queue** of **to-be-explored** nodes.

At each iteration:

- First **finish** the first element in the queue; dequeue.
- Then enqueue the out-neighbours of the dequeued element.



$Q = [0]$


Breadth First Search

Queue implementation of BFS

Maintain a **queue** of **to-be-explored** nodes.

At each iteration:

- First **finish** the first element in the queue; dequeue.
- Then enqueue the out-neighbours of the dequeued element.



$Q = [5, 6]$


Breadth First Search

Queue implementation of BFS

Maintain a **queue** of **to-be-explored** nodes.

At each iteration:

- First **finish** the first element in the queue; dequeue.
- Then enqueue the out-neighbours of the dequeued element.



$Q = [6, 2]$


Breadth First Search

Queue implementation of BFS

Maintain a **queue** of **to-be-explored** nodes.

At each iteration:

- First **finish** the first element in the queue; dequeue.
- Then enqueue the out-neighbours of the dequeued element.



$Q = [2, 4]$


Breadth First Search

Queue implementation of BFS

Maintain a **queue** of **to-be-explored** nodes.

At each iteration:

- First **finish** the first element in the queue; dequeue.
- Then enqueue the out-neighbours of the dequeued element.



$Q = [4, 3]$


Breadth First Search

Queue implementation of BFS

Maintain a **queue** of **to-be-explored** nodes.

At each iteration:

- First **finish** the first element in the queue; dequeue.
- Then enqueue the out-neighbours of the dequeued element.



$Q = [3, 9]$


Breadth First Search

Queue implementation of BFS

Maintain a **queue** of **to-be-explored** nodes.

At each iteration:

- First **finish** the first element in the queue; dequeue.
- Then enqueue the out-neighbours of the dequeued element.



$Q = [9, 7, 8]$


Breadth First Search

Queue implementation of BFS

Maintain a **queue** of **to-be-explored** nodes.

At each iteration:

- First **finish** the first element in the queue; dequeue.
- Then enqueue the out-neighbours of the dequeued element.



$Q = [7, 8]$


Breadth First Search

Queue implementation of BFS

Maintain a **queue** of **to-be-explored** nodes.

At each iteration:

- First **finish** the first element in the queue; dequeue.
- Then enqueue the out-neighbours of the dequeued element.



$Q = [8]$


Breadth First Search

Queue implementation of BFS

Maintain a **queue** of **to-be-explored** nodes.

At each iteration:

- First **finish** the first element in the queue; dequeue.
- Then enqueue the out-neighbours of the dequeued element.



$Q = []$


Breadth First Search

Queue implementation of BFS

Maintain a **queue** of **to-be-explored** nodes.

At each iteration:

- First **finish** the first element in the queue; dequeue.
- Then enqueue the out-neighbours of the dequeued element.



$Q = [1]$


Breadth First Search

Queue implementation of BFS

Maintain a **queue** of **to-be-explored** nodes.

At each iteration:

- First **finish** the first element in the queue; dequeue.
- Then enqueue the out-neighbours of the dequeued element.



$Q = []$

Breadth First Search

Maintain a field $dist(u)$ for every node u .

Procedure **bfs_explore**(G, s)

INPUT: A graph G and a starting node s

$dist(s) \leftarrow 0$

Create an empty queue Q and enqueue(Q, s)

while Q is not empty do

$u \leftarrow \text{dequeue}(Q)$ (Note u is first in Q)

 for any outgoing edge (u, v) do

 if $dist(v) = \infty$ then

 enqueue(Q, v)

$dist(v) \leftarrow dist(u) + 1$

Procedure **bfs**(G)

INPUT: A graph G and a starting node s

OUTPUT: Labeling $dist(u)$ for every u

for $u \in V$ do $dist(u) \leftarrow \infty$

for $u \in V$ do

 if $dist(u) = \infty$ then run **bfs_explore**(G, u)

Breadth First Search

Theorem.

After running `bfs_explore(G, s)`, the value of $dist(u)$ is the distance from s to u for every node u .

Breadth First Search

Theorem.

After running `bfs_explore(G, s)`, the value of $dist(u)$ is the distance from s to u for every node u .

Proof. We use the notion of **known region**: Say a node u is in the known region if $dist(u) \neq \infty$.

Breadth First Search

Theorem.

After running `bfs_explore(G, s)`, the value of $dist(u)$ is the distance from s to u for every node u .

Proof. We use the notion of **known region**: Say a node u is in the known region if $dist(u) \neq \infty$.

We need to prove that all reachable nodes u from s eventually enters the known region, and that $dist(u)$ will be correctly calculated.

Breadth First Search

Theorem.

After running `bfs_explore(G, s)`, the value of $dist(u)$ is the distance from s to u for every node u .

Proof. We use the notion of **known region**: Say a node u is in the known region if $dist(u) \neq \infty$.

We need to prove that all reachable nodes u from s eventually enters the known region, and that $dist(u)$ will be correctly calculated.

Suppose, for a **contradiction**, that u is a closest node that will **not** enter the known region.

Breadth First Search

Theorem.

After running `bfs_explore(G, s)`, the value of $\text{dist}(u)$ is the distance from s to u for every node u .

Proof. We use the notion of **known region**: Say a node u is in the known region if $\text{dist}(u) \neq \infty$.

We need to prove that all reachable nodes u from s eventually enters the known region, and that $\text{dist}(u)$ will be correctly calculated.

Suppose, for a **contradiction**, that u is a closest node that will **not** enter the known region.

Say $\text{dist}(u) = k > 0$. Then there is a node v with $\text{dist}(v) = k - 1$ and $(v, u) \in E$.

Breadth First Search

Theorem.

After running `bfs_explore(G, s)`, the value of $dist(u)$ is the distance from s to u for every node u .

Proof. We use the notion of **known region**: Say a node u is in the known region if $dist(u) \neq \infty$.

We need to prove that all reachable nodes u from s eventually enters the known region, and that $dist(u)$ will be correctly calculated.

Suppose, for a **contradiction**, that u is a closest node that will **not** enter the known region.

Say $dist(u) = k > 0$. Then there is a node v with $dist(v) = k - 1$ and $(v, u) \in E$.

Then v would enter the known region and all outgoing edges of v will be examined.

Breadth First Search

Theorem.

After running `bfs_explore(G, s)`, the value of $dist(u)$ is the distance from s to u for every node u .

Proof. We use the notion of **known region**: Say a node u is in the known region if $dist(u) \neq \infty$.

We need to prove that all reachable nodes u from s eventually enters the known region, and that $dist(u)$ will be correctly calculated.

Suppose, for a **contradiction**, that u is a closest node that will **not** enter the known region.

Say $dist(u) = k > 0$. Then there is a node v with $dist(v) = k - 1$ and $(v, u) \in E$.

Then v would enter the known region and all outgoing edges of v will be examined.

Thus u will enter the known region and $dist(u)$ will be k .

Contradiction



Breadth First Search - Complexity

Complexity Analysis

In running a BFS on G :

- Every node u in G may be enqueued and dequeued **at most once**
- Every edge may be checked **at most once**

Breadth First Search - Complexity

Complexity Analysis

In running a BFS on G :

- Every node u in G may be enqueued and dequeued **at most once**
- Every edge may be checked **at most once**

⇒ The running time of the BFS algorithm is $O(m + n)$.

DFS versus BFS

DFS: The algorithm makes **deep but narrow exploration** into the graph.

Only retreat when it runs out of new nodes.

Implementation: Use a **stack** as an auxiliary data structure. Can also be implemented **recursively**.

Running Time: $O(m + n)$.

Applications: This could be used for analyzing **reachability**, **linearizability**, **connectedness**.

DFS versus BFS

DFS: The algorithm makes **deep but narrow exploration** into the graph.

Only retreat when it runs out of new nodes.

Implementation: Use a **stack** as an auxiliary data structure. Can also be implemented **recursively**.

Running Time: $O(m + n)$.

Applications: This could be used for analyzing **reachability**, **linearizability**, **connectedness**.

BFS: The algorithm makes **shallow but broad exploration** into the graph.

Visit nodes by increasing distances.

Implementation: Could use a **queue** as an auxiliary data structure.

Running Time: $O(m + n)$

Applications: This could be used for computing **distances**.

Day 6: Distances in Graphs

Part II: Distances in Weighted Graphs

*In real applications, we need to find distances in **weighted graphs**, that are graph whose edges have integer weights.*




Weighted Graphs



*In real applications, we need to find distances in **weighted graphs**, that are graph whose edges have integer weights.*

Weighted Graph

A **weighted graph** is $G = (V, E, w)$ where (V, E) is a graph and $w : E \rightarrow \mathbb{Z}$ is a **weight function** that assigns each edge with an integer weight.



Weighted Graphs

Weighted Graph Representation


Question: How do we represent a weighted graph?

Weighted Graphs

Weighted Graph Representation

Question: How do we represent a weighted graph?

Answer: We can again have **adjacency matrix** or **adjacency list** representations.



0	2	1	∞	∞
2	0	1	2	3
1	1	0	∞	4
∞	2	∞	0	2
∞	3	4	2	0


Adjacency Matrix representation of weighted graphs

Weighted Graphs

Weighted Graph Representation

Question: How do we represent a weighted graph?

Answer: We can again have **adjacency matrix** or **adjacency list** representations.



5

B:2, C:1

A:2, C:1, D:2, E:3

A:1, B:1, E:4

B:2, E:2

B:3, D:2, C:4

Adjacency List representation of weighted graphs

Distances in Weighted Graphs

Distances Weighted Graphs

Question: Can we compute distances in a weighted graph?

Distances in Weighted Graphs

Distances Weighted Graphs

Question: Can we compute distances in a weighted graph?

Answer: Yes!


Distances in Weighted Graphs

Distances in Weighted Graphs

Question: Can we compute distances in a weighted graph?

Answer: Yes!

Subdivide edges into a sequence of unit-length edges.



Convert a **weighted graph** G into an **unweighted graph** G'


Distances in Weighted Graphs

Distances Weighted Graphs

Question: Can we compute distances in a weighted graph?

Answer: Yes!

Subdivide edges into a sequence of unit-length edges.



Convert a **weighted graph** G into an **unweighted graph** G'

However this is very inefficient, as the time complexity of BFS would depend on the sum of all weights.

Question

How do we compute distances in a weighted graph more efficiently?



Shortest Path Problem (Single-Sourced)


INPUT: a weighted graph G and a source node s

OUTPUT: the shortest path from s to all other nodes.

A “Lazy” Way for Finding Distances

Recap from Last Lecture


- We can transform a weighted graph into an unweighted one.
- The approach is very inefficient because there could be a large number of auxiliary nodes
- We don't care about the distances on these auxiliary nodes
- We could just **go to sleep** when BFS visits these auxiliary nodes
- But we need to **wake up** when BFS visits an original node



Setting Alarm Clocks

Strategy


- Maintain an **alarm clock** for each node. The time we set on an alarm clock is an **expected time** for visiting this node.
- Whenever an alarm clock goes off, wake up and check which node is reached. Then **reset** the other clocks



Setting Alarm Clocks

Strategy


- Maintain an **alarm clock** for each node. The time we set on an alarm clock is an **expected time** for visiting this node.
- Whenever an alarm clock goes off, wake up and check which node is reached. Then **reset** the other clocks



Setting Alarm Clocks

Strategy


- Maintain an **alarm clock** for each node. The time we set on an alarm clock is an **expected time** for visiting this node.
- Whenever an alarm clock goes off, wake up and check which node is reached. Then **reset** the other clocks



Setting Alarm Clocks

Strategy


- Maintain an **alarm clock** for each node. The time we set on an alarm clock is an **expected time** for visiting this node.
- Whenever an alarm clock goes off, wake up and check which node is reached. Then **reset** the other clocks



Setting Alarm Clocks

Strategy


- Maintain an **alarm clock** for each node. The time we set on an alarm clock is an **expected time** for visiting this node.
- Whenever an alarm clock goes off, wake up and check which node is reached. Then **reset** the other clocks



Setting Alarm Clocks

Strategy


- Maintain an **alarm clock** for each node. The time we set on an alarm clock is an **expected time** for visiting this node.
- Whenever an alarm clock goes off, wake up and check which node is reached. Then **reset** the other clocks



Setting Alarm Clocks

Strategy


- Maintain an **alarm clock** for each node. The time we set on an alarm clock is an **expected time** for visiting this node.
- Whenever an alarm clock goes off, wake up and check which node is reached. Then **reset** the other clocks



Setting Alarm Clocks

Strategy



- Maintain an **alarm clock** for each node. The time we set on an alarm clock is an **expected time** for visiting this node.
- Whenever an alarm clock goes off, wake up and check which node is reached. Then **reset** the other clocks



Setting Alarm Clocks

Strategy



- Maintain an **alarm clock** for each node. The time we set on an alarm clock is an **expected time** for visiting this node.
- Whenever an alarm clock goes off, wake up and check which node is reached. Then **reset** the other clocks



Setting Alarm Clocks

Strategy



- Maintain an **alarm clock** for each node. The time we set on an alarm clock is an **expected time** for visiting this node.
- Whenever an alarm clock goes off, wake up and check which node is reached. Then **reset** the other clocks



Setting Alarm Clocks

Strategy



- Maintain an **alarm clock** for each node. The time we set on an alarm clock is an **expected time** for visiting this node.
- Whenever an alarm clock goes off, wake up and check which node is reached. Then **reset** the other clocks



Setting Alarm Clocks

Strategy



- Maintain an **alarm clock** for each node. The time we set on an alarm clock is an **expected time** for visiting this node.
- Whenever an alarm clock goes off, wake up and check which node is reached. Then **reset** the other clocks



Setting Alarm Clocks

Strategy



- Maintain an **alarm clock** for each node. The time we set on an alarm clock is an **expected time** for visiting this node.
- Whenever an alarm clock goes off, wake up and check which node is reached. Then **reset** the other clocks



Setting Alarm Clocks

Strategy



- Maintain an **alarm clock** for each node. The time we set on an alarm clock is an **expected time** for visiting this node.
- Whenever an alarm clock goes off, wake up and check which node is reached. Then **reset** the other clocks



Setting Alarm Clocks

Strategy



- Maintain an **alarm clock** for each node. The time we set on an alarm clock is an **expected time** for visiting this node.
- Whenever an alarm clock goes off, wake up and check which node is reached. Then **reset** the other clocks



Setting Alarm Clocks

Strategy



- Maintain an **alarm clock** for each node. The time we set on an alarm clock is an **expected time** for visiting this node.
- Whenever an alarm clock goes off, wake up and check which node is reached. Then **reset** the other clocks



Setting Alarm Clocks

Strategy

- Maintain an **alarm clock** for each node. The time we set on an alarm clock is an **expected time** for visiting this node.
- Whenever an alarm clock goes off, wake up and check which node is reached. Then **reset** the other clocks



“Alarm Clock Algorithm”

Simulate the execution of BFS on G' starting from node s .

“Alarm Clock Algorithm”

1. Set an alarm clock for each node for time $w(s, v)$
2. Start **BFS** on G' and go to sleep
3. Repeat the following until no more alarm is left:
 4. Whenever an alarm clock goes off, wake up
 5. Pause BFS
 6. Check the current time, say T
 7. If this is u 's alarm, write $\text{dist}(u) = T$
 8. Discard this alarm clock
 9. For each neighbour v of u do:
 10. If v 's alarm is set for a time $> T + w(u, v)$,
then reset it to $T + w(u, v)$
12. Resume BFS and go back to sleep

“Alarm Clock Algorithm”

Implementing the Alarm Clocks

Question: How do we implement the **system of alarm clocks** in a program?

“Alarm Clock Algorithm”

Implementing the Alarm Clocks

Question: How do we implement the **system of alarm clocks** in a program?

We need a data structure that is able to:

- Contain a collection of integer **keys** (i.e. alarm clock times)
- **Insert(e, k)**: Add a new element e with key k to the collection
- **DeleteMin()**: Return the element with the smallest key, and remove it from the collection
- **ResetKey(e, k)**: Reset the key value of element e to a smaller value k

“Alarm Clock Algorithm”

Implementing the Alarm Clocks

Question: How do we implement the **system of alarm clocks** in a program?

We need a data structure that is able to:

- Contain a collection of integer **keys** (i.e. alarm clock times)
- **Insert(e, k)**: Add a new element e with key k to the collection
- **DeleteMin()**: Return the element with the smallest key, and remove it from the collection
- **ResetKey(e, k)**: Reset the key value of element e to a smaller value k

This is a

“Alarm Clock Algorithm”

Implementing the Alarm Clocks

Question: How do we implement the **system of alarm clocks** in a program?

We need a data structure that is able to:

- Contain a collection of integer **keys** (i.e. alarm clock times)
- **Insert(e, k)**: Add a new element e with key k to the collection
- **DeleteMin()**: Return the element with the smallest key, and remove it from the collection
- **ResetKey(e, k)**: Reset the key value of element e to a smaller value k


This is a **Priority Queue**!

Priority Queue

Priority Queues

A **priority queue** is a data structure that store a collection of $(element, key)$ pair where the *key* of an element is an integer value and allows the following operations:


- ***Insert*(e, k)**: Add a new element e with key k to the collection
- ***DeleteMin*()**: Return the element with the smallest key, and remove it from the collection
- ***ResetKey*(e, k)**: Reset the key value of element e to a smaller value k




Dijkstra's Algorithm

Implementing the “Alarm Clocks” using Priority Queue:

- Maintain a priority key **set** for each node
- Maintain a **set** of confirmed nodes




Priority Queue P




Dijkstra's Algorithm

Implementing the “Alarm Clocks” using Priority Queue:

- Maintain a priority key **set** for each node
- Maintain a **set** of confirmed nodes




Priority Queue P




Dijkstra's Algorithm

Implementing the “Alarm Clocks” using Priority Queue:

- Maintain a priority key **set** for each node
- Maintain a **set** of confirmed nodes




Priority Queue P




Dijkstra's Algorithm

Implementing the “Alarm Clocks” using Priority Queue:

- Maintain a priority key **set** for each node
- Maintain a **set** of confirmed nodes





Priority Queue P



Dijkstra's Algorithm

Implementing the “Alarm Clocks” using Priority Queue:


- Maintain a priority key **set** for each node
- Maintain a **set** of confirmed nodes




Dijkstra's Algorithm

Implementing the “Alarm Clocks” using Priority Queue:

- Maintain a priority key **set** for each node
- Maintain a **set** of confirmed nodes



Priority Queue P



Dijkstra's Algorithm

Algorithm **Dijkstra**(G, s)

INPUT: A weighted graph G , and a node s

OUTPUT: $dist(v)$ of all nodes v

$dist(s) \leftarrow 0$

Initialize a **set** $R \leftarrow \{s\}$

Initialize a **priority queue** P containing $(s, 0)$

for $u \in V, u \neq s$ **do**

$dist(u) \leftarrow \infty$

$prev(u) \leftarrow null$

$P.Insert(u, \infty)$

while P is not empty **do**

$u \leftarrow P.DeleteMin()$

 Add u to R

for $(u, v) \in E$ where $v \notin R$ **do**

if $dist(u) + w(u, v) < dist(v)$ **do**

$dist(v) \leftarrow dist(u) + w(u, v)$

$P.ResetKey(v, dist(v))$

$prev(v) \leftarrow u$

Dijkstra's Algorithm: Complexity

Note: The running time of Dijkstra's algorithm depends on the running time of priority queue implementations.

- Each node is inserted to the priority queue once
 - For each edge, we may reset the key of an element in the priority queue
 - Each node is deleted from the priority queue once
- Let $T_{in}(n)$ be the time it takes to **insert** elements to the priority queue
 - Let $T_{re}(n)$ be the time it takes to **reset key** for an element in the priority queue
 - Let $T_{de}(n)$ be the time it takes to **delete min** element from the priority queue

Dijkstra's Algorithm: Complexity

Note: The running time of Dijkstra's algorithm depends on the running time of priority queue implementations.

- Each node is inserted to the priority queue once
 - For each edge, we may reset the key of an element in the priority queue
 - Each node is deleted from the priority queue once
- Let $T_{in}(n)$ be the time it takes to **insert** elements to the priority queue
- Let $T_{re}(n)$ be the time it takes to **reset key** for an element in the priority queue
- Let $T_{de}(n)$ be the time it takes to **delete min** element from the priority queue

Running time $nT_{in}(n) + mT_{re}(n) + nT_{de}(n)$

Dijkstra's Algorithm: Complexity

Note: The running time of Dijkstra's algorithm depends on the running time of priority queue implementations.

- Each node is inserted to the priority queue once
 - For each edge, we may reset the key of an element in the priority queue
 - Each node is deleted from the priority queue once
- Let $T_{in}(n)$ be the time it takes to **insert** elements to the priority queue
- Let $T_{re}(n)$ be the time it takes to **reset key** for an element in the priority queue
- Let $T_{de}(n)$ be the time it takes to **delete min** element from the priority queue

Running time $nT_{in}(n) + mT_{re}(n) + nT_{de}(n)$

We now look at some standard priority queues.

Dijkstra's Algorithm

Different Priority Queues

Which one is better for Dijkstra's Algorithm?

- [Linked List](#):
- [Binary Heap](#):

Dijkstra's Algorithm

Different Priority Queues

Which one is better for Dijkstra's Algorithm?

- **Linked List:** $O(n^2)$
- **Binary Heap:** $O((m + n) \log n)$

Dijkstra's Algorithm

Different Priority Queues

Which one is better for Dijkstra's Algorithm?

- **Linked List:** $O(n^2)$

Preferred when there are a lot of edges, i.e., $m \geq \frac{n^2}{\log n}$

- **Binary Heap:** $O((m + n) \log n)$

Dijkstra's Algorithm

Different Priority Queues

Which one is better for Dijkstra's Algorithm?

- **Linked List:** $O(n^2)$

Preferred when there are a lot of edges, i.e., $m \geq \frac{n^2}{\log n}$

- **Binary Heap:** $O((m + n) \log n)$

Preferred when there are not a lot of edges, i.e., $m < \frac{n^2}{\log n}$

Dijkstra's Algorithm

Different Priority Queues

Which one is better for Dijkstra's Algorithm?

- **Linked List:** $O(n^2)$

Preferred when there are a lot of edges, i.e., $m \geq \frac{n^2}{\log n}$


- **Binary Heap:** $O((m + n) \log n)$

Preferred when there are not a lot of edges, i.e., $m < \frac{n^2}{\log n}$

- **Fibonacci Heap:** $O(n \log n + m)$

Better asymptotic complexity, but complicated to implement.
(Not covered in this course.)

Graph Exploration




Exploring Graphs

A **graph exploration algorithm** traverses the graph:

- The algorithm maintains a **known region** of nodes
- Each time it **picks an edge** that goes out from the known region, exploring a node outside and expanding its known region
- It stops when no more edge can be explored

The order in which new edges are picked determines the type of algorithm.

Graph Exploration



Exploring Graphs

- **DFS** picks edges based on a **stack** order
⇒ Exploration is as deep as possible
- **BFS** picks edges based on a **queue** order
⇒ Exploration is as broad as possible (hence revealing distances)
- **Dijkstra's algorithm** picks edges based on a **priority** order (on a weighted graph)
⇒ Exploration is as broad as possible (hence revealing distances)

Edsger Dijkstra's Immortality

Edsger Dijkstra 1930 - 2002

